

Python 3.6.5 |Anaconda custom (64-bit)| (default, Mar 29 2018, 13:32:41) [MSC v.1900 64 bit (AMD64)]

Type "copyright", "credits" or "license" for more information.

IPython 6.4.0 -- An enhanced Interactive Python.

Restarting kernel...

```
In [1]: runfile('E:/deep-learning_projects/pascal-voc-classification/src/main.py',  
wdir='E:/deep-learning_projects/pascal-voc-classification/src')
```

Available device = cuda

-----Epoch 1-----

100%|██████████| 358/358 [05:36<00:00, 1.07it/s]

train_loss: 2.7819, train_avg_precision:0.778

100%|██████████| 364/364 [02:00<00:00, 3.03it/s]

val_loss: 2.1037, val_avg_precision:0.868

-----Epoch 2-----

100%|██████████| 358/358 [05:38<00:00, 1.06it/s]

train_loss: 2.0955, train_avg_precision:0.861

100%|██████████| 364/364 [01:59<00:00, 3.05it/s]

val_loss: 2.1048, val_avg_precision:0.866

-----Epoch 3-----

100%|██████████| 358/358 [05:44<00:00, 1.04it/s]

train_loss: 1.8666, train_avg_precision:0.885

100%|██████████| 364/364 [02:02<00:00, 2.98it/s]

val_loss: 1.8274, val_avg_precision:0.887

-----Epoch 4-----

100%|██████████| 358/358 [05:43<00:00, 1.04it/s]

train_loss: 1.6041, train_avg_precision:0.909

100%|██████████| 364/364 [02:01<00:00, 2.99it/s]

val_loss: 1.8171, val_avg_precision:0.895

-----Epoch 5-----

100%|██████████| 358/358 [05:41<00:00, 1.05it/s]

train_loss: 1.3729, train_avg_precision:0.932

100%|██████████| 364/364 [01:56<00:00, 3.11it/s]

val_loss: 1.7336, val_avg_precision:0.906

-----Epoch 6-----

100%|██████████| 358/358 [05:43<00:00, 1.04it/s]

train_loss: 1.1237, train_avg_precision:0.952

100%|██████████| 364/364 [01:55<00:00, 3.14it/s]

val_loss: 1.7787, val_avg_precision:0.905

-----Epoch 7-----

100%|██████████| 358/358 [05:43<00:00, 1.04it/s]

train_loss: 0.9344, train_avg_precision:0.964

100%|██████████| 364/364 [01:57<00:00, 3.10it/s]

val_loss: 1.6622, val_avg_precision:0.916

-----Epoch 8-----

100%|██████████| 358/358 [05:40<00:00, 1.05it/s]

train_loss: 0.7461, train_avg_precision:0.979

100%|██████████| 364/364 [01:57<00:00, 3.09it/s]


val_loss: 1.5785, val_avg_precision:0.923


-----Epoch 9-----

```

100%|██████████| 358/358 [05:40<00:00, 1.05it/s]
train_loss: 0.6100, train_avg_precision:0.985
100%|██████████| 364/364 [01:58<00:00, 3.08it/s]
val_loss: 1.6196, val_avg_precision:0.919
-----Epoch 10-----
100%|██████████| 358/358 [05:40<00:00, 1.05it/s]
train_loss: 0.5231, train_avg_precision:0.989
100%|██████████| 364/364 [01:58<00:00, 3.08it/s]
val_loss: 1.5531, val_avg_precision:0.926
-----Epoch 11-----
100%|██████████| 358/358 [05:40<00:00, 1.05it/s]
train_loss: 0.4617, train_avg_precision:0.991
100%|██████████| 364/364 [01:58<00:00, 3.08it/s]
val_loss: 1.5592, val_avg_precision:0.926
-----Epoch 12-----
100%|██████████| 358/358 [05:39<00:00, 1.05it/s]
train_loss: 0.4193, train_avg_precision:0.993
100%|██████████| 364/364 [01:58<00:00, 3.07it/s]
val_loss: 1.5225, val_avg_precision:0.927
-----Epoch 13-----
100%|██████████| 358/358 [05:41<00:00, 1.05it/s]
train_loss: 0.4134, train_avg_precision:0.994
100%|██████████| 364/364 [01:58<00:00, 3.08it/s]
val_loss: 1.5648, val_avg_precision:0.925
-----Epoch 14-----
100%|██████████| 358/358 [05:41<00:00, 1.05it/s]
train_loss: 0.4117, train_avg_precision:0.993
100%|██████████| 364/364 [01:56<00:00, 3.14it/s]
val_loss: 1.5557, val_avg_precision:0.925
-----Epoch 15-----
100%|██████████| 358/358 [05:40<00:00, 1.05it/s]
train_loss: 0.4289, train_avg_precision:0.992
100%|██████████| 364/364 [01:56<00:00, 3.13it/s]
val_loss: 1.5356, val_avg_precision:0.928

```






Evaluating model on test set

Loading best weights

100%|██████████| 1941/1941 [12:09<00:00, 2.66it/s]

test_loss: 1.3569, test_avg_precision:0.938



```
In [2]: runfile('E:/deep-learning_projects/pascal-voc-classification/src/main.py',
wdir='E:/deep-learning_projects/pascal-voc-classification/src')
```

Reloaded modules: dataset, train, utils

Available device = cuda

-----Epoch 1-----

100%|██████████| 358/358 [05:35<00:00, 1.07it/s]

train_loss: 2.7855, train_avg_precision:0.780

100%|██████████| 364/364 [01:59<00:00, 3.04it/s]

val_loss: 2.0760, val_avg_precision:0.865

-----Epoch 2-----

100%|██████████| 358/358 [05:39<00:00, 1.05it/s]

train_loss: 2.0199, train_avg_precision:0.867

100%|██████████| 364/364 [01:57<00:00, 3.09it/s]

val_loss: 2.4253, val_avg_precision:0.844

-----Epoch 3-----

100%|██████████| 358/358 [05:40<00:00, 1.05it/s]

train_loss: 1.7363, train_avg_precision:0.899

100%|██████████| 364/364 [01:58<00:00, 3.07it/s]

val_loss: 1.9795, val_avg_precision:0.880

-----Epoch 4-----

100%|██████████| 358/358 [05:40<00:00, 1.05it/s]



train_loss: 1.4866, train_avg_precision:0.921

```

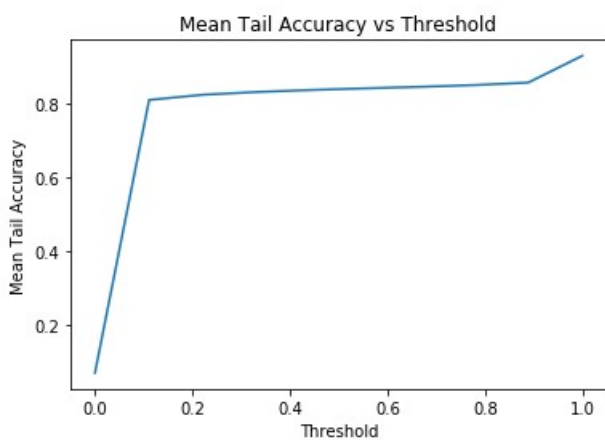
100%|██████████| 364/364 [01:59<00:00, 3.06it/s]
val_loss: 1.9375, val_avg_precision:0.887
-----Epoch 5-----
100%|██████████| 358/358 [05:43<00:00, 1.04it/s]
train_loss: 1.3005, train_avg_precision:0.938
100%|██████████| 364/364 [01:58<00:00, 3.07it/s]
val_loss: 1.8635, val_avg_precision:0.894
-----Epoch 6-----
100%|██████████| 358/358 [05:40<00:00, 1.05it/s]
train_loss: 1.0265, train_avg_precision:0.961
100%|██████████| 364/364 [01:58<00:00, 3.08it/s]
val_loss: 1.8054, val_avg_precision:0.901
-----Epoch 7-----
100%|██████████| 358/358 [05:41<00:00, 1.05it/s]
train_loss: 0.8164, train_avg_precision:0.972
100%|██████████| 364/364 [01:58<00:00, 3.06it/s]
val_loss: 1.7043, val_avg_precision:0.913
-----Epoch 8-----
100%|██████████| 358/358 [05:40<00:00, 1.05it/s]
train_loss: 0.6422, train_avg_precision:0.983
100%|██████████| 364/364 [01:58<00:00, 3.07it/s]
val_loss: 1.6481, val_avg_precision:0.920
-----Epoch 9-----
100%|██████████| 358/358 [05:39<00:00, 1.05it/s]
train_loss: 0.4944, train_avg_precision:0.990
100%|██████████| 364/364 [01:58<00:00, 3.08it/s]
val_loss: 1.6265, val_avg_precision:0.921
-----Epoch 10-----
100%|██████████| 358/358 [05:39<00:00, 1.05it/s]
train_loss: 0.4258, train_avg_precision:0.993
100%|██████████| 364/364 [01:58<00:00, 3.08it/s]
val_loss: 1.6223, val_avg_precision:0.923
-----Epoch 11-----
100%|██████████| 358/358 [05:39<00:00, 1.06it/s]
train_loss: 0.3611, train_avg_precision:0.995
100%|██████████| 364/364 [01:56<00:00, 3.11it/s]
val_loss: 1.6276, val_avg_precision:0.924
-----Epoch 12-----
100%|██████████| 358/358 [05:41<00:00, 1.05it/s]
train_loss: 0.3320, train_avg_precision:0.996
100%|██████████| 364/364 [01:57<00:00, 3.11it/s]
val_loss: 1.6022, val_avg_precision:0.925
-----Epoch 13-----
100%|██████████| 358/358 [05:41<00:00, 1.05it/s]
train_loss: 0.3249, train_avg_precision:0.996
100%|██████████| 364/364 [01:57<00:00, 3.11it/s]
val_loss: 1.6479, val_avg_precision:0.924
-----Epoch 14-----
100%|██████████| 358/358 [05:40<00:00, 1.05it/s]
train_loss: 0.3165, train_avg_precision:0.996
100%|██████████| 364/364 [01:56<00:00, 3.14it/s]
val_loss: 1.6089, val_avg_precision:0.926
-----Epoch 15-----
100%|██████████| 358/358 [05:40<00:00, 1.05it/s]
train_loss: 0.3258, train_avg_precision:0.996

```

100%|██████████| 364/364 [01:56<00:00, 3.13it/s]
val_loss: 1.6173, val_avg_precision:0.925



Evaluating model on test set
Loading best weights
100%|██████████| 1941/1941 [12:40<00:00, 2.55it/s]
test_loss: 1.4294, test_avg_precision:0.938



In [3]: